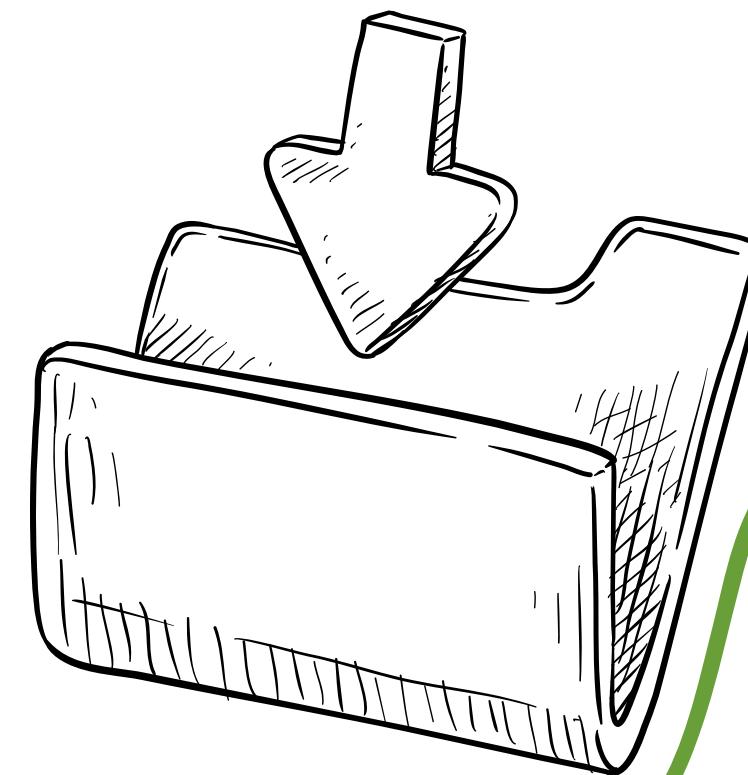
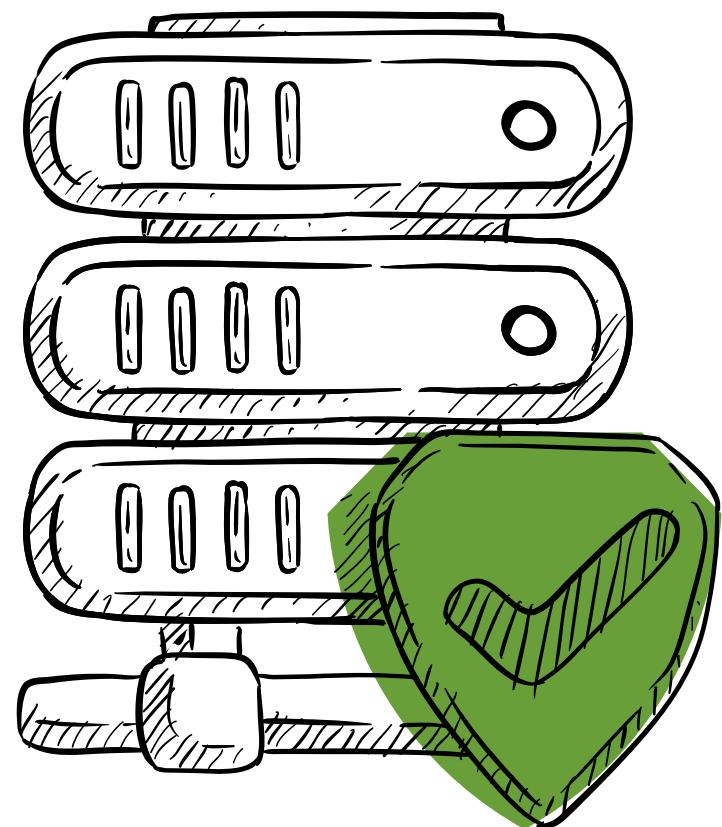


# MONGO DATABASE

## DAY03

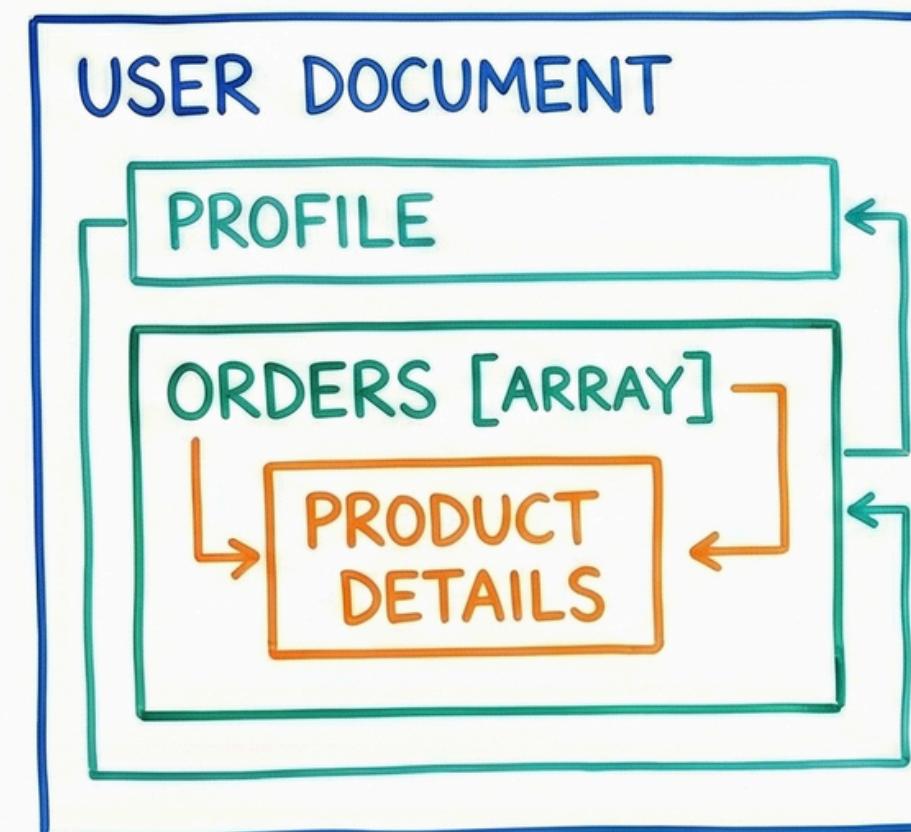
Prepared by  
Noha Shehab



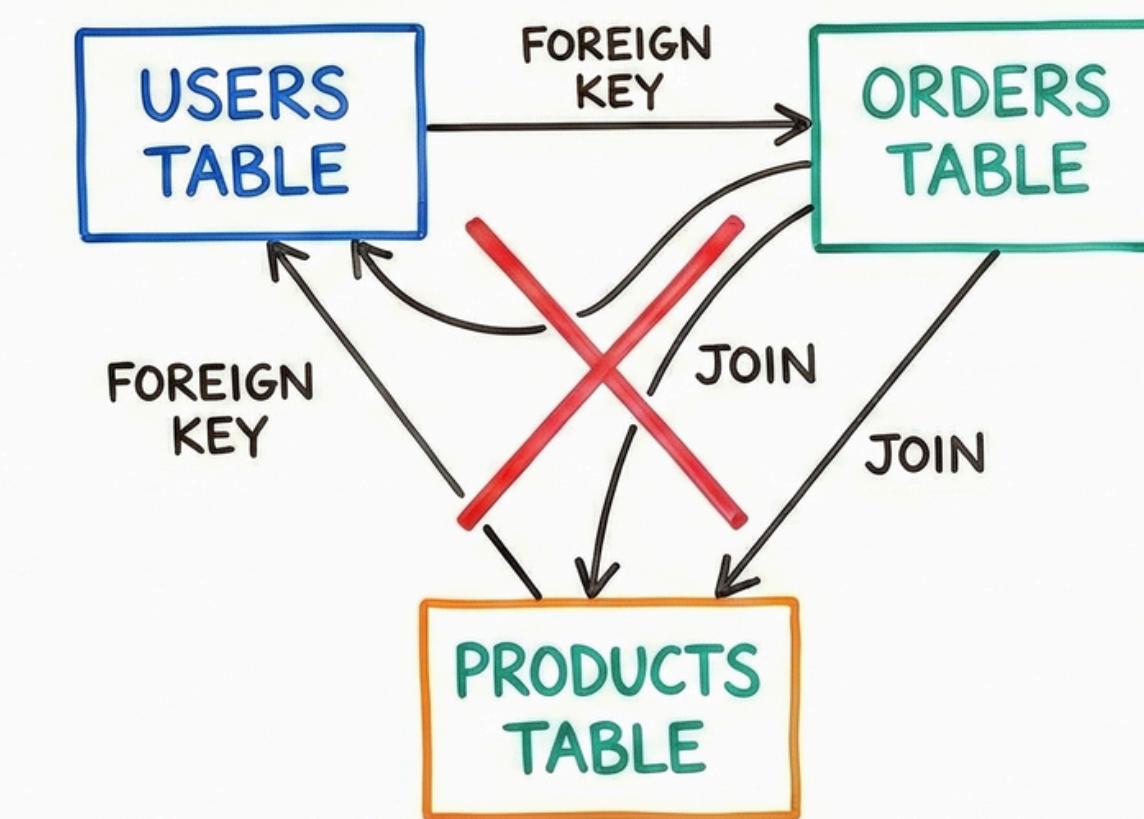
# DATA MODELING

## DATA MODELING PRINCIPLES: DOCUMENT vs. RELATIONAL (GRAPH)

### MONGODB DOCUMENT MODEL (NO GRAPH)



### RELATIONAL / GRAPH MODEL



# DATA MODELING

- Data modeling in MongoDB is fundamentally different from relational database modeling.
- In a relational database, you model based on the data's structure, **prioritizing normalization to eliminate redundancy**.
- In MongoDB, you model based on your application's access patterns.
- The primary goal is to structure your data in a way that your application can retrieve all the information it needs for a specific task in a single, efficient query.

## Core principles of data modeling

- The Golden Rule: Store Data Together That Is Accessed Together
- Design for Your Application's Query Patterns
- Prefer Embedding for Relationships
- Use Referencing for Specific Scenario
- Strategic Denormalization is Encouraged

# EMBEDDING

## Prefer Embedding for Relationships

MongoDB's document model allows you to nest documents and arrays within a parent document. This is called embedding and is the default choice for modeling relationships.

- When to Embed:
  - **"Contains" relationships:** A blog post contains comments. An order contains line items.
    - **One-to-One (1:1)** relationships: A user has one profile record.
    - **One-to-Few (1:N)** relationships: A person has a few addresses or phone numbers. The "few" side is small and bounded.

## Advantages of Embedding:

- Performance: You get all related data in a single read operation.
- Atomicity: MongoDB provides atomic writes at the document level. By embedding related data, you can update a parent and its children in a single, all-or-nothing operation without needing complex multi-document transactions.

# REFERENCING

## Use Referencing for Specific Scenarios

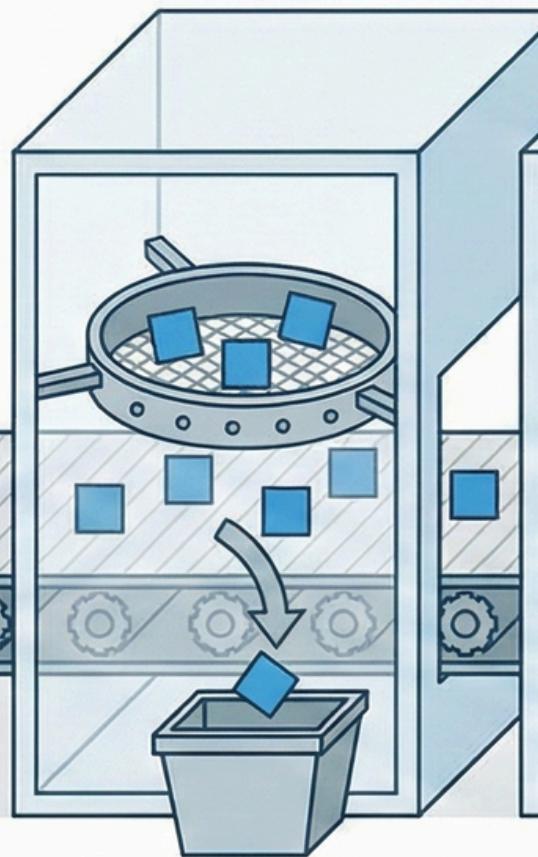
- When to Reference:
  - **One-to-Many (1:N)** with unbounded growth: A popular product could have millions of reviews. Embedding all reviews would make the product document massive and exceed MongoDB's 16MB document size limit.
  - **Many-to-Many (N:M)** relationships: Students enroll in many courses, and courses have many students. Referencing is the standard way to model this.
  - **Data is accessed independently:**
    - If you frequently need to query the "child" data without the "parent" data, referencing might be better.
  - **Data duplication is a major concern:**
    - If an entity (like a product category) is used in thousands of documents and changes frequently, embedding it would mean updating thousands of documents for a single change. Referencing avoids this.

# AGGREGATION

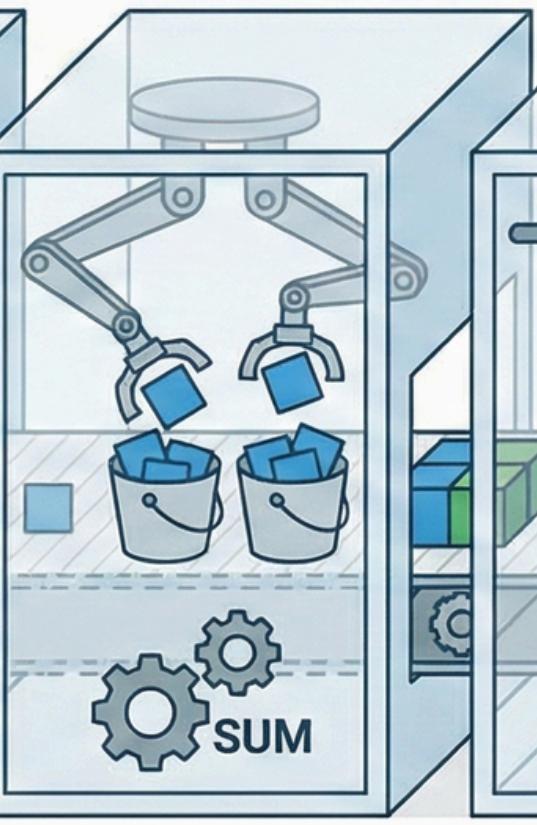
RAW DATA  
COLLECTION



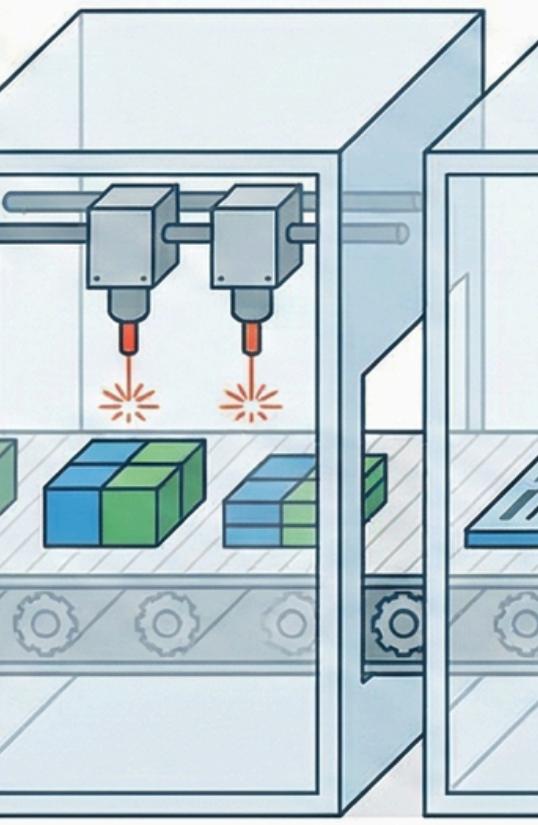
\$match  
(Filter)



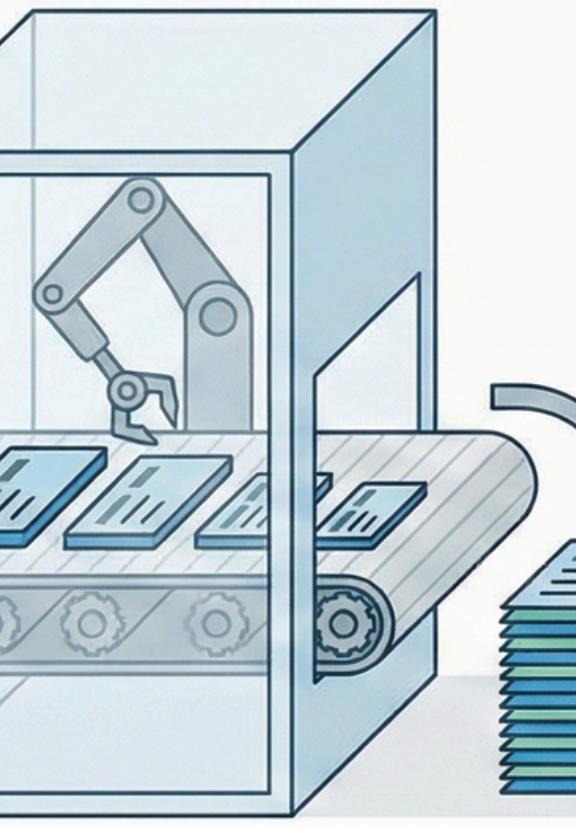
\$group  
(Aggregate)



\$project  
(Reshape)



\$sort  
(Order)



FINAL  
RESULT

\$match  
(Filter)

\$group  
(Aggregate)

\$project  
(Reshape)

\$sort  
(Order)

# AGGREGATION

## What is Aggregation?

- Aggregation is the process of taking many distinct pieces of data (documents) and processing them together to produce a computed result.
- While a standard `db.collection.find()` query is used to fetch exact copies of documents that match criteria, aggregation is used to transform, summarize, and analyze data.
- Think of it like spreadsheet operations:
  - `find()` query: Filtering a spreadsheet to show only rows where "Status = 'Active'".
  - Aggregation: Taking that filtered data and creating a **Pivot Table** to show the total sales per region, ordered by highest revenue.
- **Common use cases for aggregation:**
  - Grouping data (e.g., "How many users signed up last month by country?")
  - Performing calculations (e.g., "What is the average order value?")
  - Reshaping documents (e.g., joining data from two collections, or flattening arrays).
  - Filtering and sorting based on calculated fields.

# AGGREGATION-PIPELINE

## The Aggregation Pipeline

- The Aggregation Pipeline is the primary framework MongoDB uses to perform aggregation tasks.
- The concept is based on a data processing "pipe" or an assembly line.
  - Documents enter the pipeline and pass through a series of stages.
- An aggregation pipeline is an array [] containing stage objects {}. Every stage starts with a \$ character.

```
● ● ● aggregate.js
db.orders.aggregate([
    // Stage 1
    {
    },
    // Stage 2
    {
    },
    // Stage 3
    {
    }
])
```

# AGGREGATION-PIPELINE

## Key Principles of the Pipeline:

### 1. Sequential Processing:

- a. Documents flow through stages in the specific order you define.

### 2. Transformation:

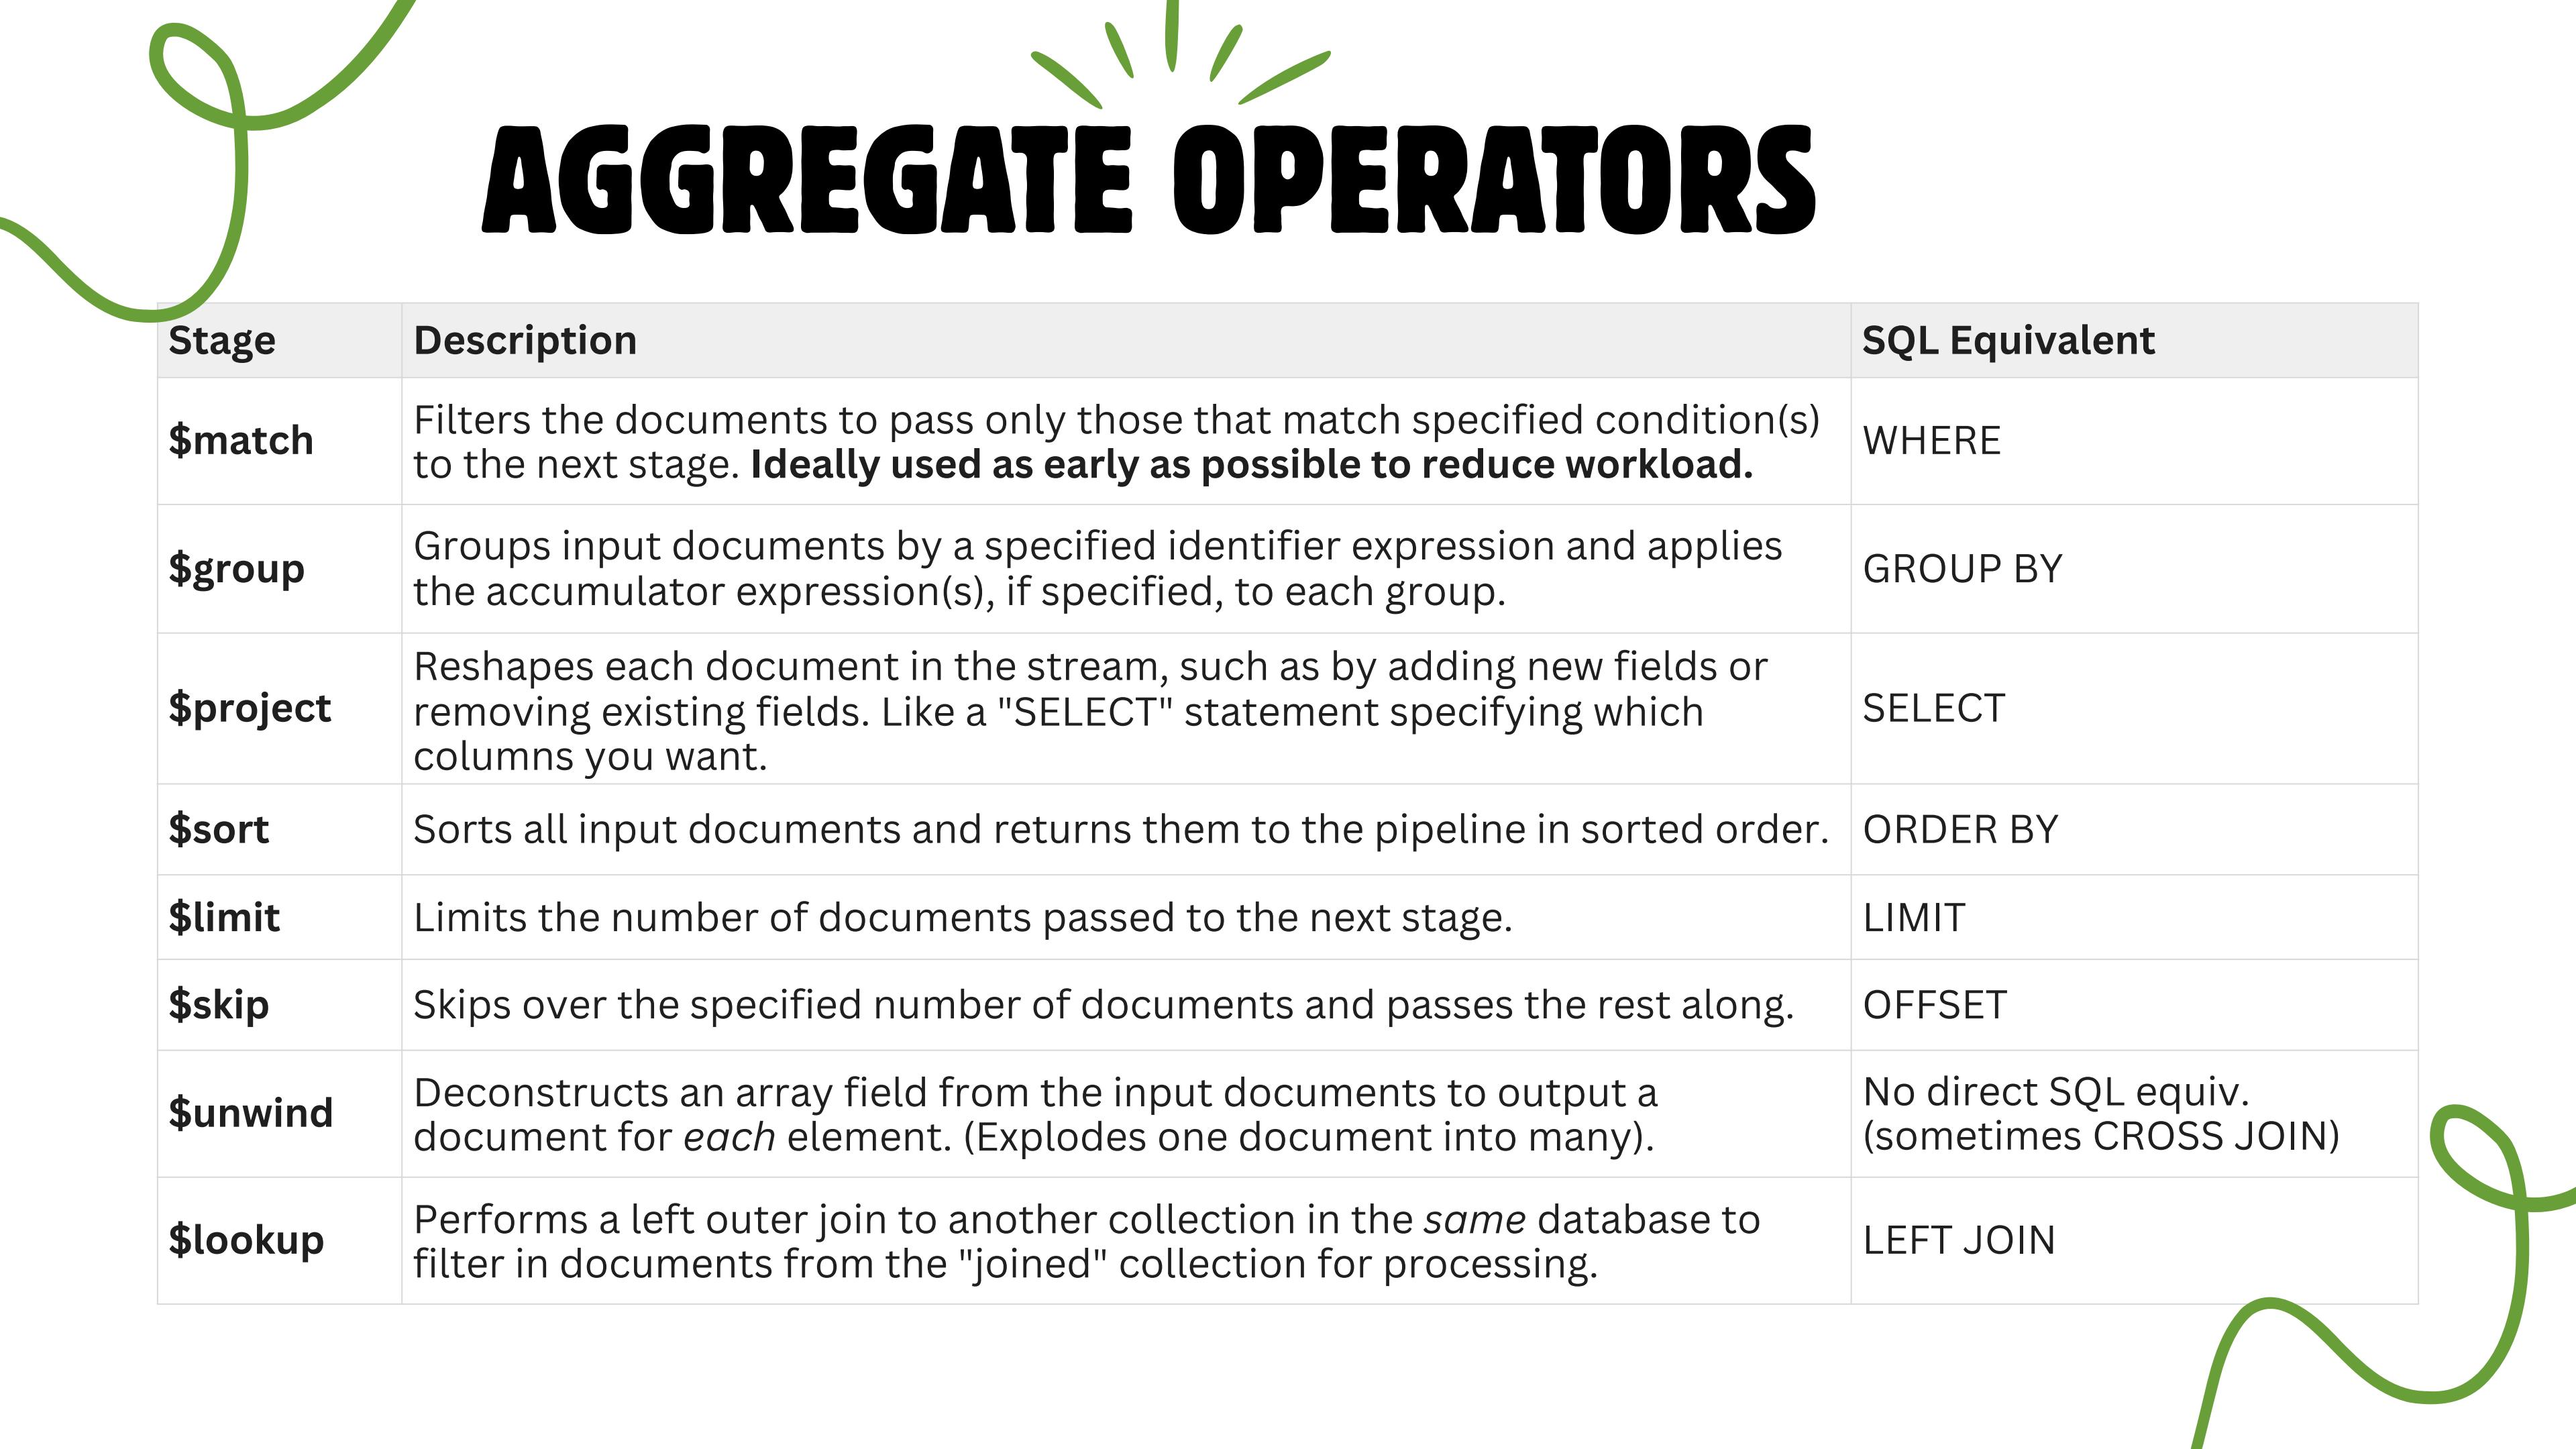
- a. Each stage performs a specific operation on the input documents and outputs transformed documents.

### 3. Input/Output:

- a. The output of one stage becomes the input for the next stage.

### 4. No Persistent Changes:

The pipeline runs in memory and returns results to the client. It does not modify the actual data in the collection (unless you specifically use an output stage like \$out or \$merge as the very last step).



# AGGREGATE OPERATORS

Stage	Description	SQL Equivalent
\$match	Filters the documents to pass only those that match specified condition(s) to the next stage. <b>Ideally used as early as possible to reduce workload.</b>	WHERE
\$group	Groups input documents by a specified identifier expression and applies the accumulator expression(s), if specified, to each group.	GROUP BY
\$project	Reshapes each document in the stream, such as by adding new fields or removing existing fields. Like a "SELECT" statement specifying which columns you want.	SELECT
\$sort	Sorts all input documents and returns them to the pipeline in sorted order.	ORDER BY
\$limit	Limits the number of documents passed to the next stage.	LIMIT
\$skip	Skips over the specified number of documents and passes the rest along.	OFFSET
\$unwind	Deconstructs an array field from the input documents to output a document for each element. (Explodes one document into many).	No direct SQL equiv. (sometimes CROSS JOIN)
\$lookup	Performs a left outer join to another collection in the <i>same</i> database to filter in documents from the "joined" collection for processing.	LEFT JOIN



# \$MATCH

- **Purpose:**
  - Filters documents.
    - It decides which documents are allowed to move to the next stage of the pipeline.
  - If an item doesn't meet the criteria (e.g., it's damaged, or it doesn't have the right security badge), it gets tossed off the line immediately. Only good items continue.

```
aggregate.js

db.instructors.aggregate([
{
  $match : {age: {$gt:21}, lastName:{$exists:true}}
} , // first stage

])
```

# \$SORT

- **Purpose:**
  - Reorders the documents in the stream based on specific fields.

```
aggregate.js

db.instructors.aggregate([
{
  $match : {age: {$gt:21}, lastName:{$exists:true}}
} ,
{
  $sort : {firstName:1 , lastName:1}
})

```

# \$PROJECT

- **Purpose:**

- Reshapes documents. You use it to select specific fields, rename fields, remove fields, or calculate entirely new fields based on existing data.

```
aggregate.js

db.instructors.aggregate([
  {
    $match: { age: { $gt: 21 }, lastName: { $exists: true } }
  },
  {
    $sort: { firstName: 1, lastName: 1 }
  },
  {

    $project: {
      fullname: { $concat: ["$firstName", ' ', "$lastName"] },
      age: 1,
      sal: "$salary",
      netsalary: { $multiply: ["$salary", .8] },
      //           _id:0 // don't do this
    }
  }
])
```

# \$out

- **Purpose:**

- Reorders the documents in the stream based on specific fields.

```
aggregate.js
db.instructors.aggregate([
  {
    $match: { age: { $gt: 21 }, lastName: { $exists: true } }
  }, /// first stage
  {
    $sort: { firstName: 1, lastName: 1 }
  },
  {
    $out: "instructors_info"
  }
])
```

# \$GROUP

- **Purpose:**

- The \$group stage collapses multiple documents into a single document based on a common value (an identifier). For each group, you can perform calculations like counting, summing, or averaging.

```
aggregate.js
db.instructors.aggregate(
  [
    {
      $match: { age: { $gt: 20 } }
    },
    {
      $group: {
        _id: "$age", // sepecify field of grouping
        total: { $sum: 1 }, // $sum → value 1 → count
        total_Ages: { $sum: "$age" }, // $sum → age
        total_salary: { $sum: "$salary" },
        min_salary: { $min: "$salary" },
        max_salary: { $max: "$salary" },
        avg_salary: { $avg: "$salary" }
      }
    },
    {
      $group: {
        _id: null,
        count: { $sum: 1 }
      }
    }
  ]
)
```

# \$LOOKUP

- **Purpose:**

- Performs a "Join" with another collection in the same database. It brings in related data from a different "table".

```
aggregate.js

db.students.aggregate([
  {
    $lookup: {
      from: "departments", // collection name I need dept name from
      localField: "department",
      foreignField: "_id",
      as: "dept_info"
    }
  } /// lookup operator → return array
])
```

# IMPORT-DATABASES

**mongorestore** is a command-line utility used to import database backups created by mongodump into a MongoDB instance.  
It recreates the databases and collections from binary BSON files.

**MONGORESTORE --DB <DATABASE\_NAME> <PATH\_TO\_BACKUP\_FOLDER>**

Flag	Description
--uri=" <connection_string>"</connection_string>	Used for remote or Atlas clusters (includes username/password).
--db <name>	Specifies which database to restore into.
--collection <name>	Specifies a single collection to restore (requires a .bson file).
--drop	<b>Important:</b> Drops every collection from the target database <i>before</i> restoring from the backup.
--nsInclude	Restores only specific namespaces (e.g., dbName.* or dbName.collectionName).
--gzip	Required if the original backup was compressed using the --gzip flag in mongodump.

# INDEXING

**Indexing** is the most effective way to improve the performance of your MongoDB queries.

- Without indexes, MongoDB must perform a collection scan, meaning it has to look through every single document in a collection to find the ones that match your query.
- With an index, MongoDB can quickly narrow down the data by looking through a sorted list of values, much like an index in the back of a textbook.
- When you create an index, MongoDB builds a specialized data structure ([a B-tree](#)) that stores a small portion of the collection's data in a searchable, sorted form.
- The Trade-off:
  - While indexes make reads (queries) much faster,
  - they slightly slow down writes (inserts/updates/deletes)
  - This because MongoDB has to update the index every time the data changes.
  - They also consume disk space and RAM

# INDEXING

## MONGODB INDEX: The Textbook Analogy

### WITHOUT INDEX

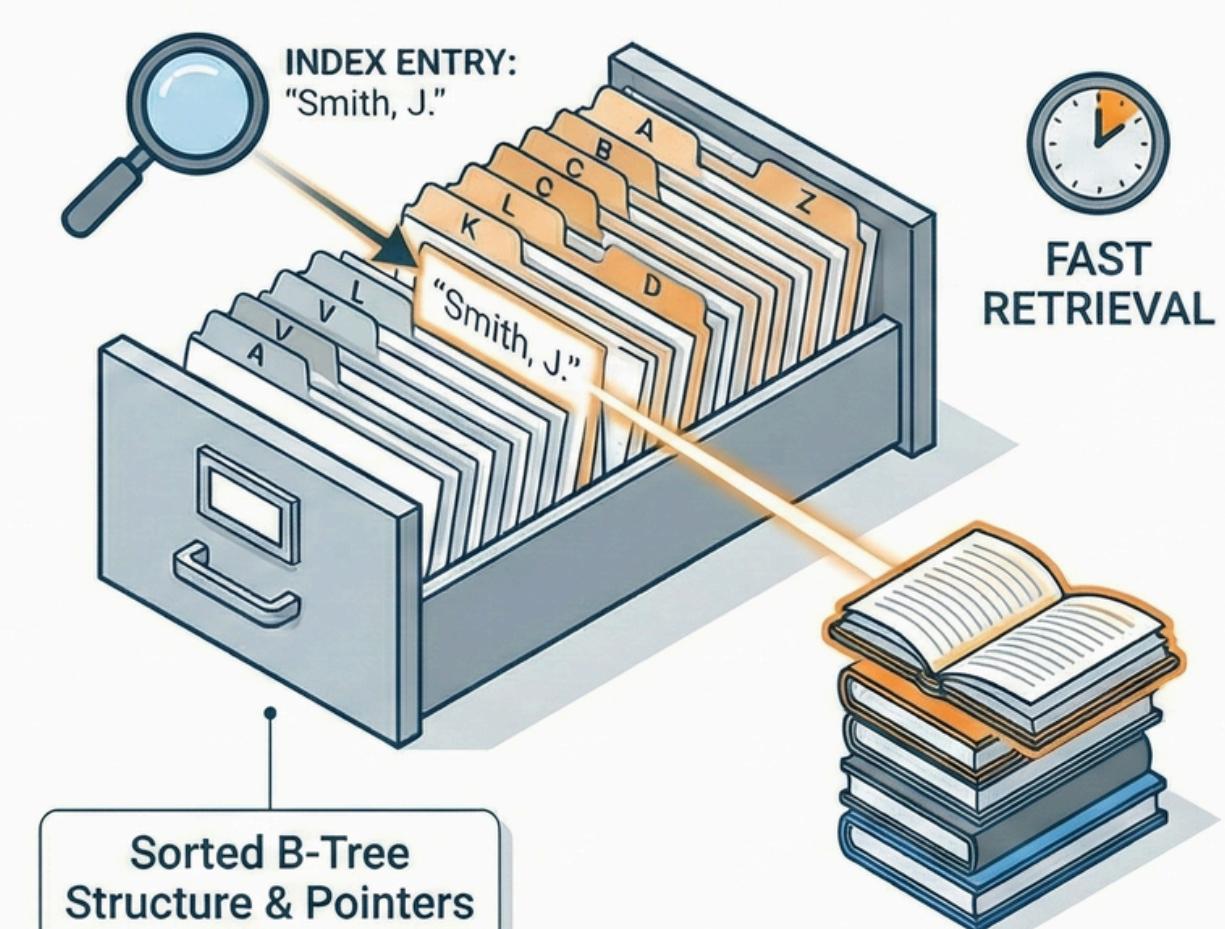
(Collection Scan)



SLOW  
SEARCH

### WITH INDEX

(Index Scan)



FAST  
RETRIEVAL

Sorted B-Tree  
Structure & Pointers

# INDEXING

Index Type	Purpose	Syntax Example
<b>Single Field</b>	The most basic index on a single key.	db.users.createIndex({ "email": 1 })
<b>Compound</b>	Indexing multiple fields together (order matters!).	db.users.createIndex({ "lastName": 1, "age": -1 })
<b>Multikey</b>	Used for indexing fields that contain arrays.	db.posts.createIndex({ "tags": 1 })
<b>Text</b>	Supports search for string content (search engine style).	db.articles.createIndex({ "content": "text" })
<b>Geospatial</b>	For location-based queries (coordinates).	db.places.createIndex({ "location": "2dsphere" })
<b>Hashed</b>	Used primarily for Sharding to ensure even data distribution.	db.users.createIndex({ "user_id": "hashed" })

The values 1 and -1 specify the sort order (Ascending and Descending). For single-field indexes, the direction doesn't usually matter, but for compound indexes, it can impact performance

# INDEXING

Feature	Without Index	With Index
Query Method	Collection Scan (Slow)	Index Scan (Fast)
Data Access	Checks every document sequentially	Jumps directly to relevant data pointers
Sorting	Requires expensive in-memory operation	Uses pre-sorted index order automatically
Write Speed	Faster	Slower (must update index on every change)
Disk/RAM Usage	Lower	Higher



restore-index.js

```
db.product.createIndex({brand_name:1 })
```

```
db.product.dropIndex({brand_name:1 })
```